

THREE TEMPORAL DIFFERENCE LEARNING ALGORITHMS

Scott O'Hara

Metrowest Developers Machine Learning Group

7/8/2020

REFERENCES

Sample-based Learning Methods (M. White and A. White), University of Alberta, Alberta Machine Intelligence Institute, Coursera.

► <https://www.coursera.org/learn/sample-based-learning-methods/>

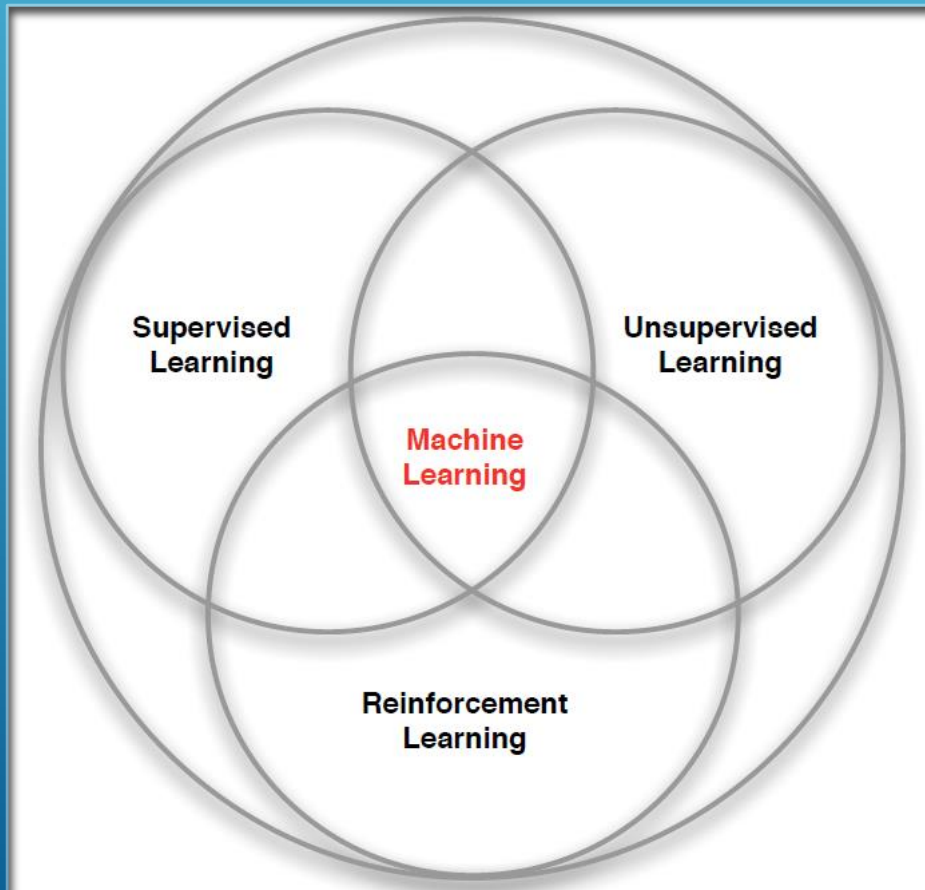
Reinforcement learning: An Introduction R. S. Sutton and A. G. Barto, Second edition. Cambridge, Massachusetts: The MIT Press, 2018.

Reinforcement Learning David Silver - University College London / Google DeepMind, 2015.

► <https://www.davidsilver.uk/teaching/>

WHAT IS REINFORCEMENT LEARNING?

“Reinforcement learning is a kind **of *unsupervised supervised learning***”
– *Rich Sutton*




Supervised Learning – Learn a function from labeled data that maps input attributes to an output.

Unsupervised Learning – Find classes, patterns or generalizations in unlabeled data.

Reinforcement Learning –An agent learns to maximize rewards while acting in an uncertain environment.

THE REINFORCEMENT LEARNING PROBLEM



- Agent must learn to act to maximize expected rewards.
- Agent knows the current state s , takes an action \mathbf{a} , receives a reward \mathbf{r} and observes the next state \mathbf{s}' .

$$S_0, A_0, R_0, S_1, A_1, R_1, S_2, A_2, R_2, \dots, S_n, A_n, R_n, S_T$$

- Agent has **no access** to the reward model $\mathbf{r(s, a, s')}$ or the transition model $\mathbf{T(s, a, s')}$.

MARKOV DECISION PROCESSES

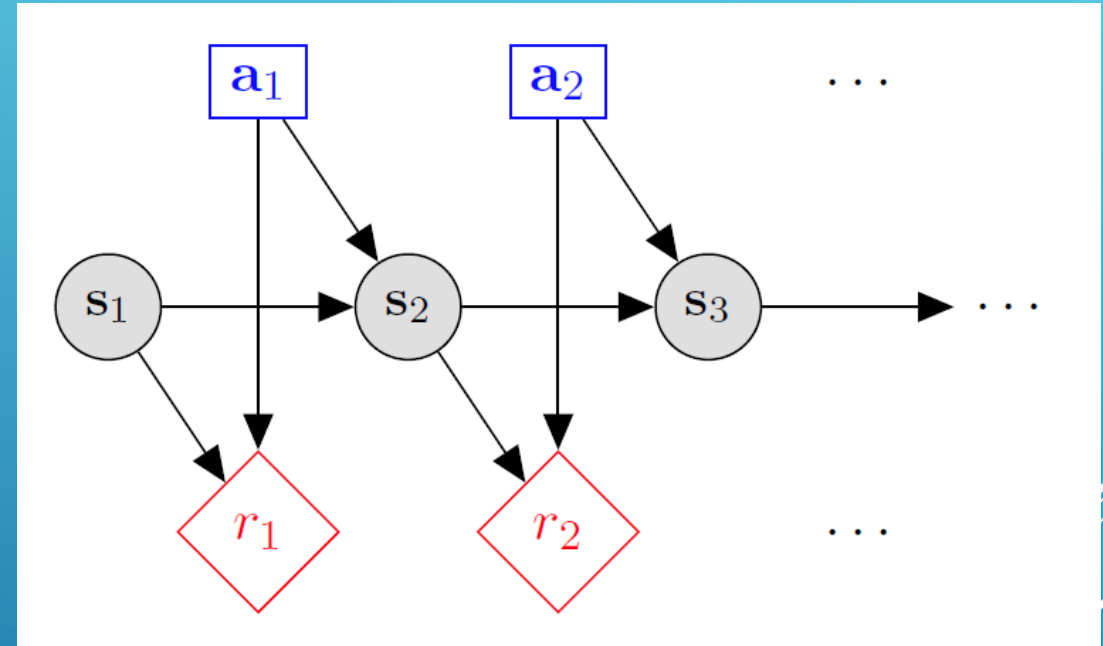
- **States:** s_1, \dots, s_n
- **Actions:** a_1, \dots, a_m
- **Reward model:**

$$R(s, a, s') \in R$$

- **Transition model:**


$$T(s, a, s') = P(s'|s, a)$$

- **Discount factor:** $\gamma \in [0, 1]$




QUANTITIES TO OPTIMIZE

- The value (utility) of a **state** s :
 $V(s)$ = expected utility starting in s and acting optimally thereafter.
- The value (utility) of a **q-state** (s,a) :
 $Q(s,a)$ = expected utility when taking action a from state s and acting optimally thereafter.
- The policy π :
 $\pi(a | s)$ = probability of action a from state s




Maze Example




- Rewards: -1 per time-step
- Actions: N, E, S, W
- States: Agent's location

Maze Example: Policy




■ Arrows represent policy $\pi(s)$ for each state s

Maze Example: Value Function



- Numbers represent value $v_{\pi}(s)$ of each state s

Maze Example: Model



- Agent may have an internal model of the environment
- Dynamics: how actions change the state
- Rewards: how much reward from each state
- The model may be imperfect

- Grid layout represents transition model $\mathcal{P}_{ss'}^a$
- Numbers represent immediate reward \mathcal{R}_s^a from each state s (same for all a)

WHAT IS TEMPORAL DIFFERENCE (TD) LEARNING?

- TD-Learning is a kind of *prediction learning* that takes advantage of the temporal structure of learning to predict.
- In prediction learning:
 - You make a prediction about what will happen next.
 - You wait to see what happens.
 - You learn by comparing what happens to what you predicted.

WHAT IS TEMPORAL DIFFERENCE (TD) LEARNING?

- TD-Learning is one of the most fundamental ideas in reinforcement learning.
- From *Reinforcement Learning: An Introduction*: “If one had to identify one idea as central and novel to reinforce learning, it would be temporal difference learning.” (page 119, Chapter 6)

Updating from a Prediction


$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

The Future
(The TD-target)

The Present
(The TD-prediction)

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, \dots, S_t, A_t, R_{t+1}, S_{t+1}$$

Updating
from a
Prediction



$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

TD Learning: Value Prediction

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in S^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal

PREDICTION VS CONTROL

- Unfortunately, the TD Value Estimation algorithm will only allow you to **predict** the value you will get from being in a given state.
- BUT, unless you have a model of the environment, it does not allow you to determine a policy to **control** the agent to maximize the value in the environment.

Update Rule:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Bellman Eq:

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

PREDICTION VS CONTROL (2)

- Fortunately, if you learn the state-action value function **Q**, you can both **predict** the value of an action at a given state **AND** you can **control** the agent to maximize the value in the environment.

Update Rule:

$$Q_k(s, a) \leftarrow Q_k(s, a) + \alpha[r + \gamma Q_k(s', a') - Q_k(s, a)]$$

Bellman Eq:

$$V^*(s) = \max_a Q^*(s, a)$$

3 UPDATE RULES, 3 ALGORITHMS

Sarsa:

$$Q_k(s, a) \leftarrow Q_k(s, a) + \alpha[r + \gamma Q_k(s', a') - Q_k(s, a)]$$

Q-Learning:

$$Q_k(s, a) \leftarrow Q_k(s, a) + \alpha \left[r + \gamma \max_{a'} Q_k(s', a') - Q_k(s, a) \right]$$

Expected Sarsa:

$$Q_k(s, a) \leftarrow Q_k(s, a) + \alpha \left[r + \gamma \sum_{a'} \pi(a'|s') Q_k(s', a') - Q_k(s, a) \right]$$

CLASSIFYING RL ALGORITHMS

- **Prediction:**
 - Predict the value of a state or a state-action pair.
 - How does the algorithm compute the value functions V and Q ?
- **Control:**
 - How does algorithm decide what to do next?
 - How is the agent's policy created and optimized?
 - Exploitation vs. exploration strategies.
- **Planning:**
 - Does the algorithm use a model of the environment?
 - How is the model created and updated?
 - How is the model exploited?

ON-POLICY / OFF-POLICY MODEL-BASED / MODEL-FREE

- **Control:**
 - **On-policy** – learn policy based on the one you are following.
 - **Off-policy** – learn policy different from the one you are following.
- **Planning:**
 - **Model-based** – use a model of the environment for prediction and control.
 - **Model-free** – learn value function or policy directly without a model i.e., the transition function and the reward function. Sarsa, Q-learning and Expected Sarsa are all model-free.

SARSA – On-policy TD Control

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Q-learning – Off-policy TD Control

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal


Expected SARSA – On-policy TD Control

Same as Q-Learning, but substitute expected state-action value for the max state-action value.

$$Q_k(s, a) \leftarrow Q_k(s, a) + \alpha \left[r + \gamma \sum_{a'} \pi(a'|s') Q_k(s', a') - Q_k(s, a) \right]$$

RL ENVIRONMENT: CLIFFWORLD


- **States:** $\langle x, y \rangle$ locations
- **Actions:** move **north**, **south**, **east** or **west**.
- **Reward model:**
 - 0 if robot moves to the goal state G where episode finishes.
 - -100 if robot moves to the cliff.
 - -1 for every other move.
- **Transition model:**
 - Robot moves deterministically in the chosen direction: **north**, **south**, **east** or **west**.
 - Robot stays put if it moves into a wall.
 - Robot transitions to the start state if it moves onto the cliff. (NOTE: episode does not finish.)



- **Discount factor:** $\gamma = 1.0$

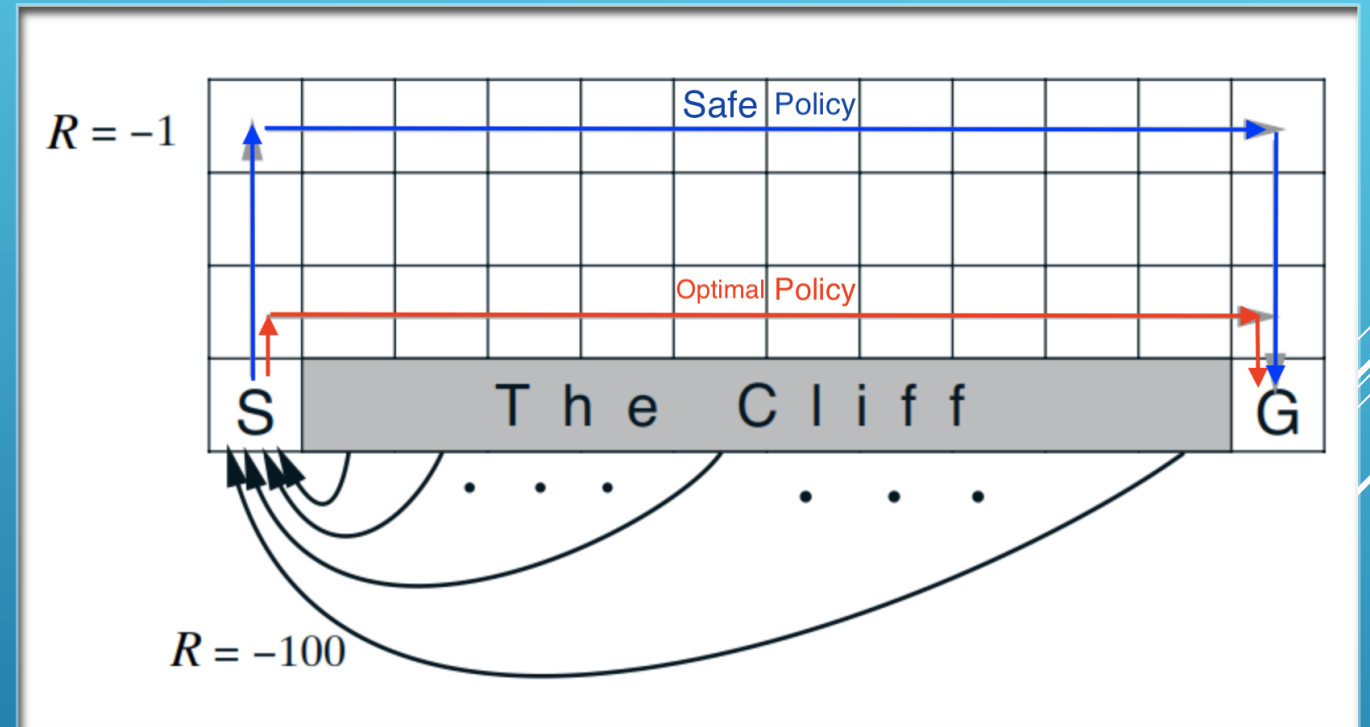
Q-LEARNING: OPTIMAL VS. “SAFE” POLICIES

- Q-learning will learn the optimal policy.
- However, q-learning must stop exploring and change to complete exploitation mode to take advantage of this.
- If q-learning continues to explore (off-policy), it will often get bad results since exploration will lead it to step over the cliff.




SARSA: OPTIMAL VS. “SAFE” POLICIES

- SARSA will learn the safe policy since it learns the policy it actually does.
- However, SARSA learns slowly since it doesn't take full advantage of the knowledge it has of state-action values.



EXPECTED SARSA: OPTIMAL VS. “SAFE” POLICIES

- Is it possible to learn an optimal policy that allows the agent to continue exploring?
- Yes! Expected SARSA, which is an on-policy algorithm can take into consideration the probability that an exploration action will take it over the cliff.
- Downside: It is more expensive to compute this policy.



CHOOSING AN ACTION: EXPLORATION VS EXPLOITATION

- How should an agent choose an action? An obvious answer is simply to follow the current policy. However, this is often not the best way to improve your model.
- **Exploit:** use your current model to maximize the expected utility now.
- **Explore:** choose an action that will help you improve your model.

EXPECTED SARSA WITH E-GREEDY METHOD

- n – number of actions
- m – number of max actions
- With probability $1 - \epsilon$:
select the action with the maximum value.

$$A_t = \operatorname{argmax} Q_t(a)$$

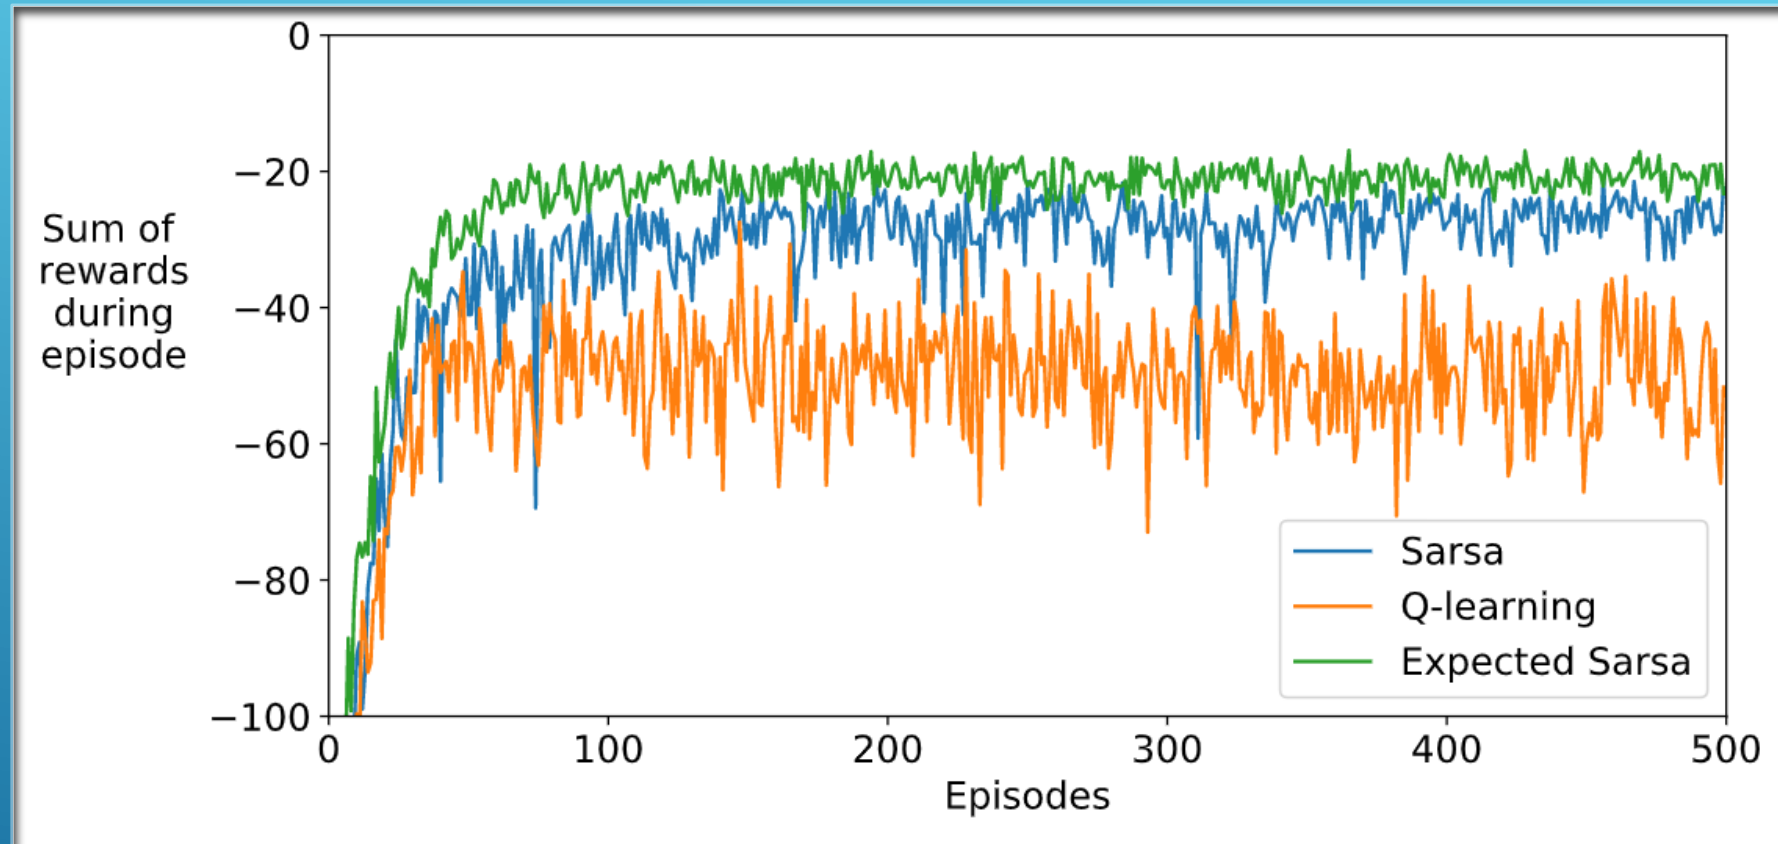
$$P(a_{\max}) = \frac{(1 - \epsilon)}{m} + \frac{\epsilon}{n}$$

- With probability ϵ :
randomly select an action from all the actions with equal probability.

$$P(a_{\text{other}}) = \frac{\epsilon}{n}$$




SARSA VS. Q-LEARNING VS. EXPECTED SARSA



- 100 runs / 500 episodes per run.
- Average the sum of rewards for each episode over 100 runs.


Q-LEARNING VS. EXPECTED SARSA

- 100 runs / 500 episodes per run.
- Average the number of visits to a state during the last 10 episodes over 100 runs.




SARSA VS. EXPECTED SARSA

- 100 runs / 500 episodes per run.
- Average the number of visits to a state during the last 10 episodes over 100 runs.



STEP-SIZE: SARSA VS. Q-LEARNING VS. EXPECTED SARSA



- 100 runs / 100 episodes per run.
- Average the sum of rewards for each episode over 100 runs for each step-size.

CONCLUSIONS

- Q-Learning will learn the optimal policy for an MDP but cannot fully exploit it unless it stops exploring.
- If q-learning continues to explore, the total value per episode will be sub-optimal.
- Expected Sarsa can find an optimal policy for a blend of exploitation and exploration.
- However, the computational overhead for Expected Sarsa is significant.
- The Sarsa algorithm can “play it safe” since it learns the policy it actually carries out.

THREE TEMPORAL DIFFERENCE LEARNING ALGORITHMS

Scott O'Hara

Metrowest Developers Machine Learning Group

7/8/2020